

Sapienza University of Rome

Master in Artificial Intelligence and Robotics  
Master in Engineering in Computer Science

## Machine Learning

A.Y. 2019/2020

Prof. L. Iocchi, F. Patrizi, V. Ntouskos

### 3. Decision Trees

L. Iocchi, F. Patrizi, V. Ntouskos

# Overview

- Decision tree representation
- ID3 learning algorithm
- Entropy, Information gain
- Inductive Bias
- Overfitting and pruning

## References

T. Mitchell. Machine Learning. Chapter 3

## Concept Learning as search

*Problem:* Given a training set  $D$  for a target function  $c$ , compute consistent hypotheses wrt  $D$ .

### *Solution approach*

- Define hypothesis space  $H$
- Implement an algorithm to search  $h \in H$  that are consistent with  $D$

# Decision Trees

Hypothesis space: set of decision trees.

## Definition

Given an instance space  $X$  formed by values coming from a set of attributes, a *decision tree* is a tree with the following characteristics:

- Each internal node tests an attribute
- Each branch denotes a value of an attribute
- Each leaf node assigns a classification value

## Example *PlayTennis*

Instances  $X = Outlook \times Temperature \times Humidity \times Wind$   
tuples of attributes

$Outlook = \{Sunny, Overcast, Rain\}$

$Temperature = \{Hot, Mild, Cold\}$

$Humidity = \{Normal, High\}$

$Wind = \{Weak, Strong\}$


Classification values:

$PlayTennis = \{Yes, No\}$

## Example *PlayTennis*: Training data

| Day | Outlook  | Temperature | Humidity | Wind   | PlayTennis |
|-----|----------|-------------|----------|--------|------------|
| D1  | Sunny    | Hot         | High     | Weak   | No         |
| D2  | Sunny    | Hot         | High     | Strong | No         |
| D3  | Overcast | Hot         | High     | Weak   | Yes        |
| D4  | Rain     | Mild        | High     | Weak   | Yes        |
| D5  | Rain     | Cool        | Normal   | Weak   | Yes        |
| D6  | Rain     | Cool        | Normal   | Strong | No         |
| D7  | Overcast | Cool        | Normal   | Strong | Yes        |
| D8  | Sunny    | Mild        | High     | Weak   | No         |
| D9  | Sunny    | Cool        | Normal   | Weak   | Yes        |
| D10 | Rain     | Mild        | Normal   | Weak   | Yes        |
| D11 | Sunny    | Mild        | Normal   | Strong | Yes        |
| D12 | Overcast | Mild        | High     | Strong | Yes        |
| D13 | Overcast | Hot         | Normal   | Weak   | Yes        |
| D14 | Rain     | Mild        | High     | Strong | No         |

## Decision Tree for *PlayTennis*



## Decision Tree for *PlayTennis*

Decision trees represent a disjunction of conjunctions of constraints on the attribute values of instances.

$$\begin{aligned} & (Outlook = Sunny \wedge Humidity = Normal) \vee \\ & (Outlook = Overcast) \vee \\ & (Outlook = Rain \wedge Wind = Weak) \end{aligned}$$

Disjunction of conjunctions of all the paths to positive (true) leaf nodes.

## Converting a Tree to Rules

A rule is generated for each path to a leaf node.

IF  $(Outlook = Sunny) \wedge (Humidity = High)$   
THEN  $PlayTennis = No$

IF  $(Outlook = Sunny) \wedge (Humidity = Normal)$   
THEN  $PlayTennis = Yes$

...

Decisions are made explicit.

## ID3 Algorithm

Input: *Examples*, *Target\_attribute*, *Attributes*

Output: *Decision Tree*

- ① Create a *Root* node for the tree
- ② **if** all *Examples* are positive, **then** return the node *Root* with label +
- ③ **if** all *Examples* are negative, **then** return the node *Root* with label –
- ④ **if** *Attributes* is empty, **then** return the node *Root* with label = most common value of *Target\_attribute* in *Examples*
- ⑤ Otherwise ...

## ID3 Algorithm

### 5. Otherwise

- $A \leftarrow$  the “best” decision attribute for *Examples*
- Assign  $A$  as decision attribute for *Root*
- For each value  $v_i$  of  $A$ 
  - add a new branch from *Root* corresponding to the test  $A = v_i$
  - $Examples_{v_i}$  = subset of *Examples* that have value  $v_i$  for  $A$
  - **if**  $Examples_{v_i}$  is empty **then** add a leaf node with label = most common value of *Target\_attribute* in *Examples*
  - **else**
    - add the tree  $ID3(Examples_{v_i}, Target\_attribute, Attributes - \{A\})$

## Which is the best attribute to choose?

*Information gain* measures how well a given attribute separates the training examples according to their target classification.


ID3 selects the attribute that has highest information gain.

Information gain measured as reduction in *entropy*.

## Entropy

- $p_{\oplus}$  is the proportion of positive examples in  $S$
- $p_{\ominus} (= 1 - p_{\oplus})$  is the proportion of negative examples in  $S$
- Entropy measures the impurity of  $S$

$$\text{Entropy}(S) \equiv -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus}$$



## Entropy example

Consider the set  $S = [9+, 5-]$  (9 positive examples, 5 negative examples)

$$Entropy(S) = -(9/14)\log_2(9/14) - (5/14)\log_2(5/14) = 0.940$$

Note: we define  $0\log_2 0 = 0$ .

Maximum entropy when  $p_{\oplus} = 0.5$ , minimum when  $p_{\oplus} = 0$  or  $1$ .

In case of multi-valued target functions ( $c$ -wise classification)

$$Entropy(S) \equiv \sum_{i=1}^c -p_i \log_2 p_i$$

## Information Gain

$Gain(S, A)$  = expected reduction in entropy of  $S$  caused by knowing the value of attribute  $A$ .

$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

$Values(A)$  : set of all possible values of  $A$

$$S_v = \{s \in S | A(s) = v\}$$



## Information Gain example

$$\text{Values}(\text{Wind}) = \{\text{Weak}, \text{Strong}\}$$

$$S = [9+, 5-]$$

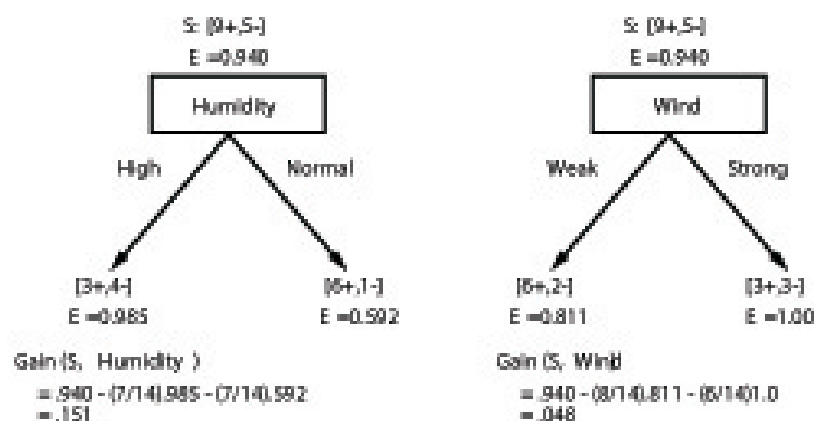
$$S_{\text{Weak}} = [6+, 2-]$$

$$S_{\text{Strong}} = [3+, 3-]$$

$$\begin{aligned} \text{Gain}(S, \text{Wind}) &= \text{Entropy}(S) - \frac{8}{14} \text{Entropy}(S_{\text{Weak}}) - \frac{6}{14} \text{Entropy}(S_{\text{Strong}}) \\ &= 0.940 - \frac{8}{14} 0.811 - \frac{6}{14} 1.00 \\ &= 0.048 \end{aligned}$$

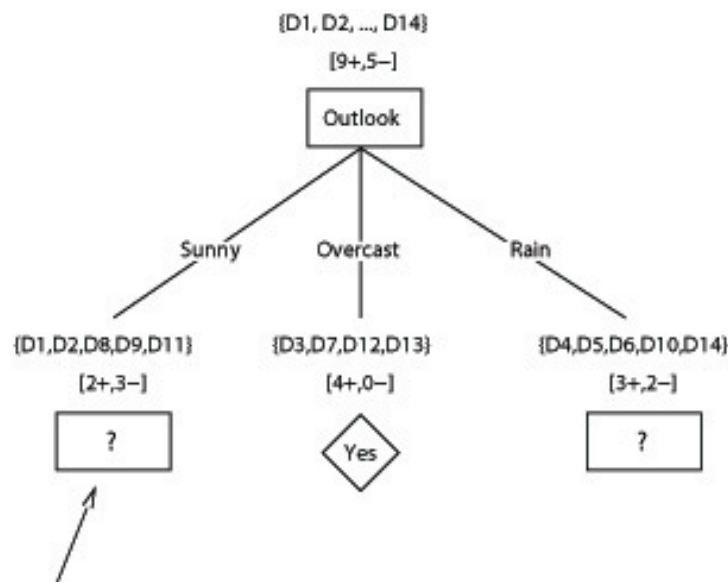
## Selecting the Next Attribute

Which attribute is the best classifier?



$$\begin{aligned} \text{Gain}(S, \text{Outlook}) &= 0.246 \\ \text{Gain}(S, \text{Humidity}) &= 0.151 \\ \text{Gain}(S, \text{Wind}) &= 0.048 \\ \text{Gain}(S, \text{Temperature}) &= 0.029 \end{aligned}$$

## Selecting the Next Attribute




$S_{\text{sunny}} = \{D1, D2, D8, D9, D11\}$

$\text{Gain}(S_{\text{sunny}}, \text{Humidity}) = .970 - (3/5) 0.0 - (2/5) 0.0 = .970$


$\text{Gain}(S_{\text{sunny}}, \text{Temperature}) = .970 - (2/5) 0.0 - (2/5) 1.0 - (1/5) 0.0 = .570$

$\text{Gain}(S_{\text{sunny}}, \text{Wind}) = .970 - (2/5) 1.0 - (3/5) .918 = .019$

## Decision Tree for *PlayTennis*



# Hypothesis Space Search by ID3



# Hypothesis Space Search by ID3

- Hypothesis space is complete (target concept is there!)
- Outputs a single hypothesis (cannot determine how many DTs are consistent!)
- No back tracking (local minima!)
- Statistically-based search choices (robust to noisy data!)
- Uses all the training examples at each step (not incremental!)

# Issues in Decision Tree Learning

- Determining how deeply to grow the DT
- Handling continuous attributes
- Choosing appropriate attribute selection measures
- Handling training data with missing attribute values
- Handling attributes with different costs

## Overfitting in Decision Trees

Consider a new data set  $D' = D \cup d_{15}$  adding a training example:

$$d_{15} = \langle \textit{Sunny}, \textit{Hot}, \textit{Normal}, \textit{Strong}, \textit{PlayTennis} = \textit{No} \rangle$$

ID3 will generate a different tree  $T'$

Note:  $T$  is consistent with  $D$  and  $T'$  is consistent with  $D'$  (i.e., accuracy = 100%)

Is  $T'$  in general a better solution for our learning problem?


# Overfitting in Decision Trees

Let's consider a more complex problem with  $|D| \gg 15$  and containing noisy data and two decision trees  $T$  and  $T'$  obtained with different configuration of an ID3-like algorithm.

$$accuracy_D(T') > accuracy_D(T)$$

Is  $T'$  in general a better solution for our learning problem?

## Overfitting in Decision Tree Learning



# Avoiding Overfitting

How can we avoid overfitting?

- stop growing when data split not statistically significant
- grow full tree, then post-prune

To determine the correct tree size

- use a separate set of examples (distinct from the training examples) to evaluate the utility of post-pruning
- apply a statistical test to estimate accuracy of a tree on the entire data distribution
- using an explicit measure of the complexity for encoding the examples and the decision trees.

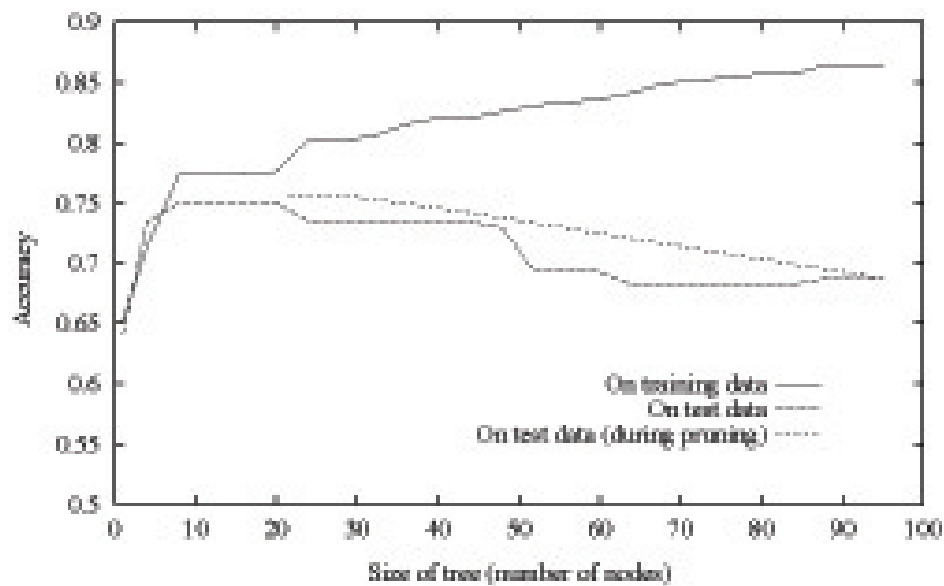
## Reduced-Error Pruning

Split data into *training* and *validation* set

Do until further pruning is harmful (decreases accuracy):

- 1 Evaluate impact on *validation* set of pruning each possible node (remove all the subtree and assign the most common classification)
- 2 Greedily remove the one that most improves *validation* set accuracy

## Effect of Reduced-Error Pruning



## Remarks on Reduced-Error Pruning

- it produces smallest version of most accurate subtree (removing sub-trees added due to coincidental irregularities).
- When data set is limited, reducing the set of training examples (used as validation examples) can give bad results.

## Rule Post-Pruning (C4.5)

- ① Infer the decision tree allowing for overfitting
- ② Convert the learned tree into an equivalent set of rules
- ③ Prune (generalize) each rule independently of others
- ④ Sort final rules into desired sequence for use

## Continuous Valued Attributes

Create a discrete attribute to test continuous variables

- $Temperature = 82.5$
- $(Temperature > 72.3) = t, f$

|                     |    |    |     |     |     |    |
|---------------------|----|----|-----|-----|-----|----|
| <i>Temperature:</i> | 40 | 48 | 60  | 72  | 80  | 90 |
| <i>PlayTennis:</i>  | No | No | Yes | Yes | Yes | No |



## Attributes with Many Values

Problem:

- If attribute has many values, *Gain* will select it
- Imagine using *Date = Jun\_3\_1996* as attribute

One approach: use *GainRatio* instead

$$\text{GainRatio}(S, A) \equiv \frac{\text{Gain}(S, A)}{\text{SplitInformation}(S, A)}$$

$$\text{SplitInformation}(S, A) \equiv - \sum_{i=1}^c \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$$

where  $S_i$  is subset of  $S$  for which  $A$  has value  $v_i$

## Attributes with Costs

Consider

- medical diagnosis, *BloodTest* has cost \$150
- robotics, *Width\_from\_1ft* has cost 23 sec.

How to learn a consistent tree with low expected cost?

Replace gain by

- Tan and Schlimmer (1990)

$$\frac{\text{Gain}^2(S, A)}{\text{Cost}(A)}$$

- Nunez (1988) ( $w \in [0, 1]$  determines importance of cost)

$$\frac{2^{\text{Gain}(S, A)} - 1}{(\text{Cost}(A) + 1)^w}$$

## Unknown Attribute Values

What if some examples missing values of  $A$ ?

Use training example anyway, sort through tree

- If node  $n$  tests  $A$ , assign most common value of  $A$  among other examples sorted to node  $n$
- assign most common value of  $A$  among other examples with same target value
- assign probability  $p_i$  to each possible value  $v_i$  of  $A$ 
  - assign fraction  $p_i$  of example to each descendant in tree

Classify new examples in same fashion

## Other algorithms based on Decision Trees

*Random Forest*: ensemble method that generates a set of decision trees with some random criteria and integrates their values into a final result.

Random criteria: 1) random subsets of data (bagging), 2) random subset of attributes (feature selection), ...

Integration of results: majority vote (most common class returned by all the trees).

Random Forests are less sensitive to overfitting.

# Summary

- Decision Trees can represent classification function by **making decisions explicit**
- Learning as search in the hypothesis space with heuristics based on information gain
- Statistical method (some robustness to noisy data)
- Overfitting and pruning
- Used as basis of randomized ensemble methods